Gorre Dinesh Chandan Reddy

+91 7416767453 | gorredinesh21@gmail.com | linkedin.com/in/gorre-dinesh-chandan-reddy | portfolio | github.com/gorredinesh21 | Mumbai, India

PROFESSIONAL SUMMARY

Associate AI/ML Engineer with strong software engineering fundamentals shipping production-ready, scalable AI services on a Databricks / PySpark / cloud stack. At Reliance Industries, builds production Generative AI applying RAG, LangChain/LangGraph, and vector retrieval (FAISS, ChromaDB) - including a RAG framework that scales Databricks Genie past its 30-table limit to 120+ tables across 9 spaces - plus near-real-time PySpark pipelines on Azure/ADF. IIT (ISM) Dhanbad CSE graduate with hands-on depth in REST/FastAPI services, deep learning, model training and quantitative evaluation, and MLOps.

CORE COMPETENCIES

Generative AI & RAG | LangChain / LangGraph | Vector Databases (FAISS, ChromaDB) | REST / FastAPI Services | Databricks & PySpark / SparkSQL | Azure / AWS Cloud Platforms | Deep Learning & Model Evaluation | MLOps & Data Pipelines

WORK EXPERIENCE

Reliance Industries Ltd.

Aug 2025 - Present

Graduate Engineer Trainee (Data / AI Engineer) - Mumbai, India

- **Unified Genie Architecture (Production GenAI / RAG)** - Designed and operated a production-ready, scalable RAG framework with LangChain that overcomes Databricks Genie's 30-table limitation by organizing 120+ tables across 9 Genie spaces, applying Generative AI to enterprise data.
- **Intelligent Query Routing Engine (Vector Retrieval)** - Built semantic retrieval pipelines over table/column metadata embeddings to dynamically route prompts to the relevant Genie space.
- **Talk-to-Genie Service (REST / FastAPI)** - Built an external web application exposing Databricks Genie through OAuth-secured REST APIs, enabling secure conversational access outside Databricks end to end.
- **Near Real-Time Data Pipeline (Databricks / PySpark / Azure)** - Built ADF ingestion pipelines on the Databricks/ADLS Gen2 stack supporting 15-minute refresh cycles with Landing to Bronze to Silver transformations.

PROJECTS

GAN-Based Data Augmentation for Imbalanced Datasets

Model Training & Quantitative Evaluation

Built an end-to-end deep-learning augmentation framework using a custom GAN and SMOTE to address class imbalance on a 10,000,000+ row credit-card dataset. Designed Generator/Discriminator networks with Xavier init, LeakyReLU, and BatchNorm for stable tabular GAN training. Trained Random Forest, XGBoost, LightGBM, and Gradient Boosting on real vs. real+augmented data with 5-fold cross-validation, baselining performance with quantitative metrics - recall improved 0.91 to 0.99 and false negatives cut to just 9.

Python, PyTorch, GANs, XGBoost, LightGBM, scikit-learn

CareerOS - Multi-Agent LLM Platform

Generative AI / LangGraph / FastAPI

Built an end-to-end Generative AI platform extracting job postings from LinkedIn, Indeed, and Telegram and generating tailored resumes per job description. Orchestrated a multi-agent LLM workflow with LangChain/LangGraph that scores fit probability and produces customised resumes, exposed through an interactive dashboard with analytics and downloadable output.

Python, Flask, FastAPI, SQLite, LangChain, LangGraph

Tinder AI Coach - Multimodal GenAI Service

RAG / FastAPI / Multimodal

Built a FastAPI service that standardizes and optimizes a user's dating profile and writes it back to the live account. Integrated Google Gemini for multimodal photo analysis and curated an India-context knowledge base that grounds advice via retrieval, with Hugging Face and Ollama model options.

Python, FastAPI, LangChain, Google Gemini, Hugging Face, Ollama

EDUCATION

B.Tech, Computer Science and Engineering - Indian Institute of Technology (ISM), Dhanbad

Dec 2021 - May 2025

Relevant coursework: Data Structures & Algorithms, OOP, DBMS, Software Engineering, Computer Networks, Deep Learning. Secured AIR 2903 in JEE Advanced; Top 200 of 18,500+ teams in the Amazon ML Challenge.

SKILLS

Generative AI: RAG, LangChain, LangGraph, Vector Databases (FAISS, ChromaDB), Agentic AI, MCP, Fine-tuning, Hugging Face, Ollama

Data & Platform Engineering: Databricks, PySpark / SparkSQL, ADF, ADLS Gen2, Apache Spark, ETL, Data Warehousing, Delta Lake / Unity Catalog (Databricks governance), AWS, Azure

Software Engineering & APIs: REST APIs, FastAPI, Flask, OAuth, MLOps / model lifecycle

ML / DL: Machine Learning, Deep Learning, NLP, CNN, LSTMs, BERT, Transformers, Transfer Learning, GANs, Computer Vision, TensorFlow, PyTorch, Keras, model training & quantitative evaluation

Languages / Databases: Python, SQL, PostgreSQL, C++, C